


High-level Intro to Benchmarking and Profiling

Håvard Anda Estensen, Vipps

Agenda

- Why profile your system?
- Benchmarking vs. Profiling
- Benchmark functions
- Profile systems
- Dos and Don'ts
- Do you even need this?!

Why Profile our System?

- We want our system to be fast 
 - Amazon: +100ms => -1% revenue
 - Our program is only part of the system
- Don't guess what the performance should be, measure it!

**“The First Rule of Program Optimization:
Don't do it”**

**“The Second Rule of Program Optimization
(for experts only):
Don't do it yet.”**

– Michael A. Jackson

Benchmarking vs. Profiling

- Benchmarking measure the execution time of a function
- Profiling is about finding performance bottlenecks

Benchmark Functions

Big O notation

do

you speak it?

```
$ time wc -w sherlock.txt
1095695 sherlock.txt
wc -w sherlock.txt
    0.02s user
    0.00s system
    95% cpu
    0.027 total
```



```
func ConcatStrings(words []string) string {  
    var str string  
    for _, word := range words {  
        str += word  
    }  
    return str  
}
```

```
func BenchmarkConcatString(b *testing.B) {  
    words := make([]string, 100)  
    for i := range words {  
        words[i] = RandStringBytes(length)  
    }  
    b.Run("concat string", func(b *testing.B) {  
        for i := 0; i < b.N; i++ {  
            ConcatStrings(words)  
        }  
    })  
}
```

```
$ go test -bench=ConcatStrings
BenchmarkConcatStrings/concat_strings-12
103142
10387 ns/op
```

```
b.Run("concat string", func(b *testing.B) {  
    b.ReportAllocs()  
    ...  
})
```

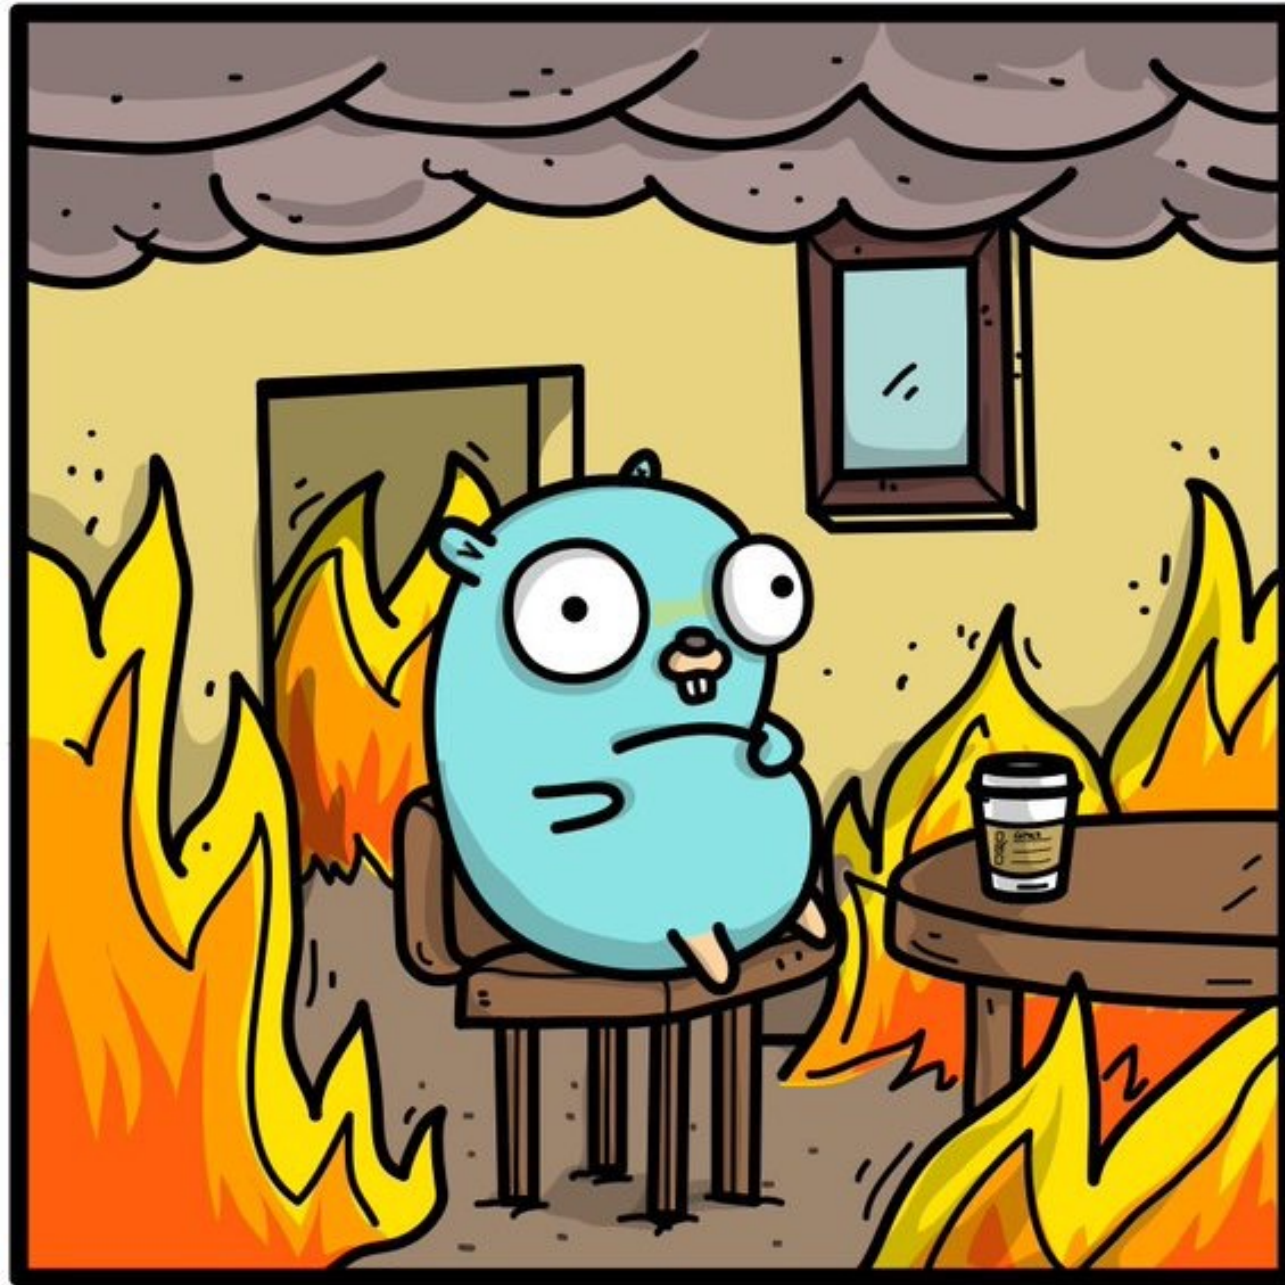
```
$ go test -bench=ConcatStrings  
BenchmarkConcatStrings/concat_strings-12  
103142  
10387 ns/op  
53488 B/op  
99 allocs/op
```

```
b.Run("concat string", func(b *testing.B) {  
    b.ReportAllocs()  
    ...  
})
```

```
$ go test -bench=ConcatStrings  
BenchmarkConcatStrings/concat_strings-12  
110922  
10566 ns/op  
53488 B/op  
99 allocs/op
```

```
$ go get golang.org/x/perf/cmd/benchstat
$ go test -bench=BenchmarkConcatString -count=8 >> old.txt
$ benchstat old.txt
```

name	time/op
ConcatString/concat_string-12	11.3µs ± 9%



**But why are we concatenating
100 strings? Is this a realistic
use-case for our program?**

**Is this our performance
bottleneck?**

CPU Profile

- Interrupt program every 10ms and record collection of stack traces
- Save profile to disk
- If the program runs for less than 10ms the file will be empty 🙄

```
import _ "net/http/pprof"
```

```
or
```

```
import "runtime/pprof"
```

```
or
```

```
"github.com/pkg/profile"
```

```
func main() {  
    defer profile.Start(profile.CPUProfile, profile.ProfilePath("."))  
    ...  
}
```

```
"github.com/pkg/profile"
```

```
func main() {  
    defer profile.Start(profile.CPUProfile, profile.ProfilePath("."))  
    ...  
}
```

```
$ go tool pprof -http=:8080 mem.pprof
```

```
import _ "net/http/pprof"
```

```
$ go tool pprof -http=:8080 http://localhost:8080/debug/pprof/profile
```

Demo: leftpad

You pay for memory allocation more than once. The first time when you allocate it. And then every time the garbage collector runs.

“Reduce/Reuse/Recycle” ❤️

–@bboreham

Stack vs. Heap

- Each thread has its own stack
- Threads shares heap
- Coordination is expensive
- GC reclaim stuff on the heap
- Conclusion: Try to allocate on the stack

```
defer profile.Start(profile.TraceProfile, profile.ProfilePath(".")).Stop()
```

```
$ go tool trace trace.out
```

Dos and Don'ts

- Use idle and dedicated hw
- Throttle CPU and turn off power saving
- Run benchmarks multiple times
- Use ~~complicated math~~ benchstat to measure ~~variance~~ noise
- Regression benchmarks are useful
- Don't do it!

Who Should Optimize?

- Do you have enough RAM?
- Can you avoid creating a distributed system?
- Does your hourly report take two hours to run?
- Every optimization comes at a cost of time and most often complexity

“Don’t optimize, but also don’t be dumb”

–Damian Gryski

Conclusion

- Don't spend time prematurely optimizing
- Profile with Idle CPU on same machine
- Preferably multiple times with perflock to avoid 🔥
- Heap allocation is expensive, stack allocation is cheap
- Spend time on improving your bottleneck

Resources

- Code for the presentation
<https://github.com/estensen/go-perf>
- Optimizing Go Code Without a Blindfold
https://www.youtube.com/watch?v=oE_vm7KeV_E
- Two Go Programs, Three Different Profiling Techniques
<https://www.youtube.com/watch?v=nok0aYiGiYA>
- Profiling and optimizing Go web applications
<https://artem.krylysov.com/blog/2017/03/13/profiling-and-optimizing-go-web-applications/>

@estensen on Keybase

